



Akhila ✓

@akki.ai

50 LLM questions interviewer
expects you to know:



Akhila ✓

@akki.ai

1. What are Large Language Models (LLMs)?
2. How are LLMs different from traditional language models?
3. What are Foundation Models?
4. How is GPT-4 different from GPT-3 in capabilities and applications?
5. What is tokenization, and why is it important in LLMs?
6. What role do embeddings play in LLMs?
7. How do LLMs handle out-of-vocabulary (OOV) words?
8. What are Sequence-to-Sequence (Seq2Seq) Models?
9. What is the difference between an encoder and a decoder?
10. How do autoregressive models differ from masked models?



Akhila ✓

@akki.ai

11. What is masked language modeling, and how does it help pretraining?
12. What is Next Sentence Prediction (NSP)?
13. How does the Transformer architecture overcome challenges of traditional Seq2Seq models?
14. What are positional encodings in LLMs?
15. What is attention in Transformer models?
16. What is Multi-head attention?
17. How are attention scores computed in transformers?
18. What is the softmax function and why is it used in attention?
19. How is dot product used in self-attention?
20. What is beam search, and how does it differ from greedy decoding?



Akhila ✓
@akki.ai

**We've curated a complete AI
Interview Preparation Bundle with
22 courses & 1100+ questions.**

Check out the bundle link in caption.



Akhila ✓

@akki.ai

21. Explain the concept of temperature in LLM text generation.
22. Difference between top-k sampling and top-p sampling?
23. How does Adaptive Softmax speed up LLMs?
24. What is cross-entropy loss and why is it used in language models?
25. How are gradients computed with respect to embeddings?
26. What is the role of the Jacobian matrix in backpropagation?
27. What is the chain rule and how is it used in deep learning?
28. What is ReLU and why is it important?
29. What is the vanishing gradient problem, and how do Transformers solve it?
30. What are eigenvalues and eigenvectors (in dimensionality reduction)?



Akhila ✓

@akki.ai

31. How is KL divergence used in evaluating LLMs?
32. What is LoRA and QLoRA?
33. How can catastrophic forgetting be mitigated in LLMs?
34. How does Parameter-Efficient Fine-Tuning (PEFT) prevent catastrophic forgetting?
35. What is model distillation, and how is it applied to LLMs?
36. What is overfitting, and how can it be prevented?
37. What are Generative and Discriminative models?
38. Difference between Discriminative AI and Generative AI?
39. How does prompt engineering influence the output of LLMs?
40. What is Chain-of-Thought (CoT) prompting?



Akhila ✓

@akki.ai

41. What are the key steps in the Retrieval-Augmented Generation (RAG) pipeline?
42. How does knowledge graph integration help LLMs?
43. How does Gemini's architecture improve training efficiency and stability compared to other multimodal LLMs?
44. How does Mixture of Experts (MoE) improve scalability?
45. What is zero-shot learning in LLMs?
46. What is few-shot learning in LLMs and why is it useful?
47. What is a context window?
48. What is a hyperparameter?
49. An LLM generates offensive or incorrect outputs. How do you handle it?
50. What are common challenges of using LLMs?



Akhila ✓
@akki.ai

**We've curated a complete AI
Interview Preparation Bundle with
22 courses & 1100+ questions.**

Check out the bundle link in caption.



Akhila ✓

@akki.ai

Thanks for reading!

Save this thread and share with a friend.